



Logistic Regression For Binomial Classification

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


Problem:

You have data of students that studied for certain hours and their pass or fail status.

Here,
0 -> Fail
1-> Pass

Design a machine learning model that can determine pass/fail status of a student who scored 3.4.

<u>Study Hours</u>	<u>Pass/Fail</u>
1.0	0
2.5	0
3.0	0
4.5	1
5.0	1
6.0	1



Logistic Regression For Binomial Classification

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


This problem solved by python and scikit-learn

<u>Study Hours</u> <u>(X)</u>	<u>Pass/Fail</u> <u>(y)</u>
1.0	0
2.5	0
3.0	0
4.5	1
5.0	1
6.0	1

```
from sklearn.linear_model import LogisticRegression
import numpy as np

# Data (Study Hours and Pass/Fail)
X = np.array([[1], [2.5], [3], [4.5], [5], [6]]) # Feature
y = np.array([0, 0, 0, 1, 1, 1]) # Target

# Create and train the model
model = LogisticRegression()
model.fit(X, y)

# Predict probabilities for new students
hours = np.array([[3.4]])
probs = model.predict_proba(hours)

# Display the probabilities of passing (class 1)
for h, p in zip(hours.flatten(), probs[:, 1]):
    print(f"Study Hours: {h}, Probability of Passing: {p:.2f}")
```

Study Hours: 3.4, Probability of Passing: 0.41



Logistic Regression For Binomial Classification

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


Sometime we have to predict whether a data point belongs to certain category. Category can be

1. Binomial: There can be only two possible types of categories

Example: Customer is going to **buy or not buy**;

Student either **Pass or Fail**;

email is **Spam or Not Spam**;

Patient has **Cancer or no-cancer**

2. Multinomial: There can be 3 or more possible types of categories

Example: Image is either **cat, dog, or sheep**;

Iris flower can be either **setosa, versicolor or virginica**

We need a model that can classify data points in classes/categories:

This is where **Logistic Regression** come.

Logistic Regression is a model that predicts **probability** (a number between 0 and 1) and then applies a **threshold** (usually 0.5) to determine which class the data belongs to.

Logistic Regression For Binomial Classification

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



To perform binomial classification, we use **sigmoid function**:

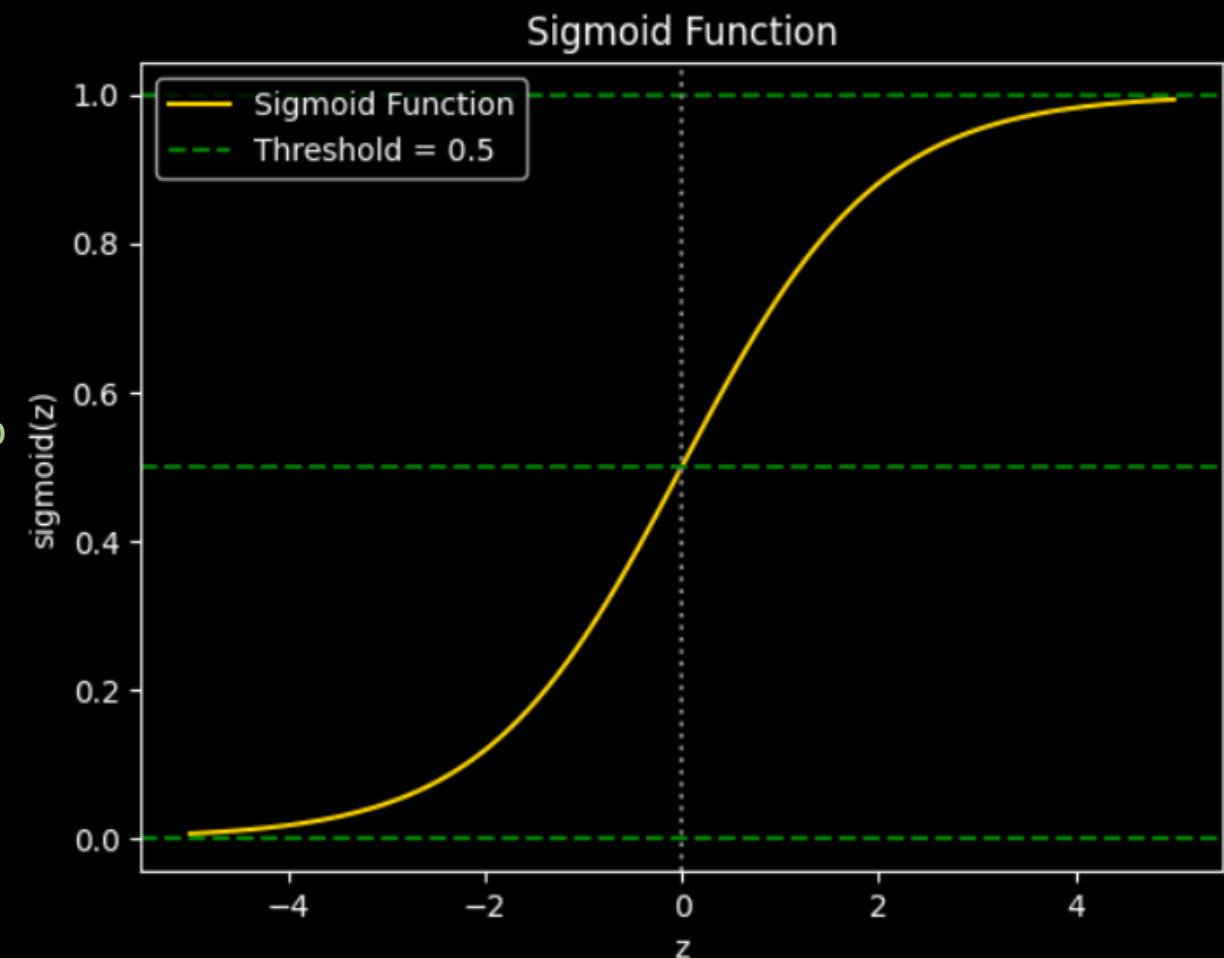
$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Important features of sigmoid function:

- Maps any real value of z to a range between **0 and 1**.
- Often used in **logistic regression** to convert linear score z into probability. i.e. The probability of obtaining z
- It is used to classify objects in 2 classes
- e is Euler Number, $2.71828\dots$ It is an irrational number

Example: Calculate probability of $z=2$

$$\begin{aligned}\sigma(2) &= \frac{1}{1 + e^{-2}} \\ &= \frac{1}{1 + 0.1353} \approx 0.8808\end{aligned}$$





Logistic Regression For Binomial Classification

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


$$x \longrightarrow z = b_0 + b_1 x \longrightarrow \sigma(z) = \frac{1}{1 + e^{-z}}$$

Step-by-Step Intuition

1. Logistic regression first calculates a **linear score** (In our example, x is StudyHours):

$$z = b_0 + b_1 \times \text{StudyHours}$$

The bias b_0 and weight b_1 are calculated during training from data .

2. Then applies the **sigmoid function** to convert z into a probability:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

The output is a probability of **Pass (1)** and is between 0 and 1.

3. If $P(\text{Pass}) > 0.5 \rightarrow \text{predict Pass (1)}$

If $P(\text{Pass}) \leq 0.5 \rightarrow \text{predict Fail (0)}$



Logistic Regression For Binomial Classification

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$


$$x \longrightarrow z = b_0 + b_1 x \longrightarrow \sigma(z) = \frac{1}{1 + e^{-z}}$$

Example: Assuming that from training data, we have calculated bias $b_0 = -4.35$ and weight $b_1 = 1.17$. and we want to find whether the student who scored 3.4, did he/she pass.

Step-by-Step:

1. Logistic regression first calculates a **linear score** (x is StudyHours):

$$z = b_0 + b_1 \times \text{StudyHours}$$

$$z = -4.35 + 1.17 \times 3.4 = -0.372$$

2. Then applies the **sigmoid function** to convert z into a probability:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\sigma(-0.372) = \frac{1}{1 + e^{0.372}}$$

$$\sigma(-0.372) = 0.41 = 41\%$$

This is the probability of **Pass (1)**.

3. Here, $P(\text{Pass}) \leq 0.5 \rightarrow \text{predict Fail (0)}$



Recap On Work Flow Of Logistic Regression


$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Machine Learning Model Before Training

$b_0 = ?$ and $b_1 = ?$

Machine Learning Model After Training

Training Data →

Calculates
 $b_0 = -4.35$ and
 $b_1 = 1.17$

Feed $x = 3.4$ →

$x \rightarrow z \rightarrow \sigma(z)$
 $3.4 \rightarrow -0.372 \rightarrow 0.41$

Logistic Regression with Many Features

Let's say you have many features: study hours, sleep time before the exam, previous score, etc

When you have multiple features:

$$X = (x_1, x_2, x_3, \dots, x_n) = (\text{study hours}, \text{sleep time}, \text{previous score})$$

Logistic regression models the probability that the output $y = 1$.

Step1: For **many features**, the model first computes:

$$z = b_0 + b_1x_1 + b_2x_2 + b_3x_3 + \dots + b_nx_n$$

where,

b_0 = bias/intercept

b_1, b_2, \dots, b_n are called weights. These are calculated during training from the training data.

x_1, x_2, \dots, x_n are called features

Step2: Then applies the **sigmoid function** to convert z into a probability:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Step3: If $P(1) > 0.5 \rightarrow \text{predict class 1}$

If $P(1) \leq 0.5 \rightarrow \text{predict class 0}$

Logistic Regression with Many Features

Python code that calculates the probability for data that has 3 features:

X1	X2	X3	y
2.5	1.3	0.5	1
1.0	3.5	2.2	0
3.2	0.9	1.4	1
4.1	2.2	0.1	1
0.9	4.5	3.1	0
3.9	1.2	0.4	1

```
# logistic regression for many features
import numpy as np
import pandas as pd
from sklearn.linear_model import LogisticRegression

# Features: X1, X2, X3
X = np.array([
    [2.5, 1.3, 0.5],
    [1.0, 3.5, 2.2],
    [3.2, 0.9, 1.4],
    [4.1, 2.2, 0.1],
    [0.9, 4.5, 3.1],
    [3.9, 1.2, 0.4]
])

# Target values (binary)
y = np.array([1, 0, 1, 1, 0, 1])

# Train logistic regression
model = LogisticRegression()
model.fit(X, y)

# Predict probability for a new sample
x_new = np.array([[3.0, 1.5, 0.7]])

prob = model.predict_proba(x_new)
print("Probability y=1:", prob[0][1])
```

Probability y=1: 0.9006659215352736

ADVANCED





Logistic Regression



Assumptions of Logistic Regression

1. Independent observations:

There should be no correlation or dependence between the input samples.

2. Linearity relationship between independent variables and log odds:

The model assumes a linear relationship between the independent variables and the log odds of the dependent variable which means the predictors affect the log odds in a linear way.

3. No outliers:

The dataset should not contain extreme outliers as they can distort the estimation of the logistic regression coefficients.

4. Large sample size:

It requires a sufficiently large sample size to produce reliable and stable results.

Logistic Regression: Find the bias b_0 and weight b_1

Goal of Logistic Regression:

We're modeling the probability that an outcome $y = 1$ (e.g., student passes) given input x (e.g., study hours):

$$P(y = 1|x) = \frac{1}{1 + e^{-(b_0 + b_1 x)}}$$

where

- b_0 = intercept (bias term)
- b_1 = slope (how much study hours influence probability)

Step 1: The Model's Prediction

For any student with study hours x_i :

$$\hat{y}_i = P(y_i = 1|x_i) = \frac{1}{1 + e^{-(b_0 + b_1 x_i)}}$$

and

$$1 - \hat{y}_i = P(y_i = 0|x_i)$$



Logistic Regression: Find the bias b_0 and weight b_1



Step 2: Likelihood Function

We want the model to assign **high probability to the actual outcomes** in the training data.

For all n data points:

$$L(b_0, b_1) = \prod_{i=1}^n [\hat{y}_i]^{y_i} [1 - \hat{y}_i]^{(1-y_i)}$$

- If $y_i = 1$, the term becomes \hat{y}_i
- If $y_i = 0$, the term becomes $1 - \hat{y}_i$

This is called the **likelihood function** — it represents how likely our parameters b_0, b_1 make the data we observe.



Logistic Regression: Find the bias b_0 and weight b_1



Step 3: Log-Likelihood Function

It's easier to work with logs (to turn the product into a sum):

$$\ell(b_0, b_1) = \sum_{i=1}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

This is the **log-likelihood function** — it's what we want to **maximize** to find the best b_0 and b_1 .

Step 4: Optimization (Finding b_0 and b_1)

To find optimal parameters:

$$\text{maximize } \ell(b_0, b_1)$$

or equivalently,

$$\text{minimize } -\ell(b_0, b_1)$$

which is called the **log loss** (or **cross-entropy loss**).

Or minimize the **Binary cross-entropy cost function**



$$J = -\frac{1}{n} \ell(b_0, b_1)$$

Logistic Regression: Find the bias b_0 and weight b_1

Step 5: Gradient Descent Intuition

We use **iterative optimization methods** — typically **Gradient Descent** or **Newton-Raphson**, to solve above problem

We start with random guesses for b_0 , b_1 , and iteratively update:

$$b_j := b_j + \alpha \frac{\partial \ell}{\partial b_j}$$

$$b_j := b_j + \alpha \sum_{i=1}^n (y_i - \hat{y}_i) x_{ij}$$

where

- α = learning rate (step size)
- $(y_i - \hat{y}_i)$ = prediction error

Over many iterations, the parameters move in the direction that **increases the likelihood** (i.e., reduces log loss).



Logistic Regression: Find the bias b_0 and weight b_1



Question: Why we do not use derivative to solve ?

In logistic regression:

$$\hat{y}_i = \frac{1}{1 + e^{-(b_0 + b_1 x_i)}}$$

We maximize the **log-likelihood**:

$$\ell(b_0, b_1) = \sum_i [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

Now, take derivatives:

$$\frac{\partial \ell}{\partial b_0} = \sum_i (y_i - \hat{y}_i)$$

$$\frac{\partial \ell}{\partial b_1} = \sum_i (y_i - \hat{y}_i) x_i$$

If we try to set these equal to zero, we get:

$$\sum_i (y_i - \hat{y}_i) = 0, \quad \sum_i (y_i - \hat{y}_i) x_i = 0$$

But \hat{y}_i itself is a nonlinear function of b_0 and b_1 (because of the exponential).

So these equations are **nonlinear** in the parameters — there's no algebraic way to isolate b_0, b_1 . They're coupled inside an exponential and a fraction.

Hence, **no closed-form solution**.

Logistic Regression: Find the bias b_0 and weight b_1

Step-by-step numeric example: Below is an example that shows how b_0 and b_1 change using **gradient ascent on the log-likelihood** (equivalently: gradient updates that *increase* the log-likelihood).

I use a tiny toy dataset of **3 students** so I can show each arithmetic step on the paper.

- Data:

$$x = [1, 2, 3] \text{ (study hours)}$$

$$y = [0, 0, 1] \text{ (0 = fail, 1 = pass)}$$

- Model:

$$\hat{y}_i = \sigma(z_i), \quad z_i = b_0 + b_1 x_i$$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- Gradients of log-likelihood (for all data points):

$$g_{b_0} = \sum_{i=1}^n (y_i - \hat{y}_i), \quad g_{b_1} = \sum_{i=1}^n (y_i - \hat{y}_i) x_i$$

- Gradient ascent update:

$$b_j \leftarrow b_j + \alpha g_{b_j}$$

(We use **gradient ascent** to *maximize* log-likelihood. If you prefer minimizing negative log-likelihood, update signs flip.)

- Hyperparameters / initial values:

$$b_0^{(0)} = 0, \quad b_1^{(0)} = 0, \quad \alpha = 0.1$$

Logistic Regression: Find the bias b_0 and weight b_1

Iteration 1 (initial $b_0 = 0, b_1 = 0$)

1. Compute linear scores z_i :

$$z = [0 + 0 \cdot 1, 0 + 0 \cdot 2, 0 + 0 \cdot 3] = [0, 0, 0]$$

2. Sigmoid (predicted probabilities) $\hat{y}_i = \sigma(z_i)$:

$$\hat{y} = [\sigma(0), \sigma(0), \sigma(0)] = [0.5, 0.5, 0.5]$$

3. Errors $e_i = y_i - \hat{y}_i$:

$$e = [0 - 0.5, 0 - 0.5, 1 - 0.5] = [-0.5, -0.5, 0.5]$$

4. Gradients:

$$g_{b_0} = \sum e_i = -0.5$$

$$g_{b_1} = \sum e_i x_i = (-0.5) \cdot 1 + (-0.5) \cdot 2 + (0.5) \cdot 3 = -0.5 - 1 + 1.5 = 0.0$$

5. Update parameters:

$$b_0 \leftarrow 0 + 0.1 \times (-0.5) = -0.05$$

$$b_1 \leftarrow 0 + 0.1 \times 0.0 = 0.0$$

After iteration 1: $b_0 = -0.05, b_1 = 0.0$



Logistic Regression: Find the bias b_0 and weight b_1



Iteration 2 (start $b_0 = -0.05, b_1 = 0$)

1. Linear scores:

$$z = [-0.05, -0.05, -0.05]$$

2. Sigmoid:

$$\sigma(-0.05) \approx 0.4875026035157896$$

So

$$\hat{y} \approx [0.4875026035, 0.4875026035, 0.4875026035]$$

3. Errors:

$$e \approx [-0.4875026035, -0.4875026035, 0.5124973965]$$

4. Gradients:

$$g_{b_0} = \sum e_i \approx -0.4625078105473689$$

$$g_{b_1} = \sum e_i x_i \approx (-0.4875026) \cdot 1 + (-0.4875026) \cdot 2 + (0.5124974) \cdot 3 \approx 0.07498437890526222$$

5. Update:

$$b_0 \leftarrow -0.05 + 0.1 \times (-0.4625078105473689) \approx -0.09625078105473689$$

$$b_1 \leftarrow 0 + 0.1 \times 0.07498437890526222 \approx 0.007498437890526222$$

After iteration 2: $b_0 \approx -0.0962507811, b_1 \approx 0.0074984379$



Logistic Regression: Find the bias b_0 and weight b_1



Iteration 3 (start $b_0 \approx -0.09625078$, $b_1 \approx 0.00749844$)

1. Linear scores:

$$z_1 \approx -0.09625078 + 0.00749844 \cdot 1 \approx -0.08875234316421067$$

$$z_2 \approx -0.09625078 + 0.00749844 \cdot 2 \approx -0.08125390527368445$$

$$z_3 \approx -0.09625078 + 0.00749844 \cdot 3 \approx -0.07375546738315822$$

2. Sigmoids:

$$\hat{y} \approx [0.4778264673349967, 0.479697692439044, 0.481569487361292]$$

3. Errors:

$$e \approx [-0.4778264673, -0.4796976924, 0.5184305126]$$

4. Gradients:

$$g_{b_0} = \sum e_i \approx -0.4390936471353327$$

$$g_{b_1} = \sum e_i x_i \approx (-0.4778264673) \cdot 1 + (-0.4796976924) \cdot 2 + (0.5184305126) \cdot 3 \approx 0.11806968570303944$$

5. Update:

$$b_0 \leftarrow -0.09625078105473689 + 0.1 \times (-0.4390936471353327) \approx -0.14016014576827016$$

$$b_1 \leftarrow 0.007498437890526222 + 0.1 \times 0.11806968570303944 \approx 0.019305406460830166$$

After iteration 3: $b_0 \approx -0.1401601458$, $b_1 \approx 0.01930540646$

EXTRA



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



Logistic Regression



$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

Quick Summary Formula Sheet

$$\hat{y}_i = \frac{1}{1 + e^{-(b_0 + b_1 x_i)}}$$

$$\ell(b_0, b_1) = \sum_i [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

$$b_j := b_j + \alpha \sum_i (y_i - \hat{y}_i) x_{ij}$$



4. Cost Function (Loss Function)

To train the model, we minimize the **Binary Cross-Entropy Loss** (also called **Log Loss**):

Given m training examples $\{(x^{(i)}, y^{(i)})\}_{i=1}^m$, where $y^{(i)} \in \{0, 1\}$:

- ◆ **Cost for one sample:**

$$\mathcal{L}(h_w(x^{(i)}), y^{(i)}) = - \left[y^{(i)} \log(h_w(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_w(x^{(i)})) \right]$$

- ◆ **Cost over all samples:**

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(h_w(x^{(i)}), y^{(i)})$$



Cost function explained on next slide

How does logistic regression works? (p2.1)

What is the **Cost Function** in Logistic Regression?

The cost function tells us **how well** our model's predictions match the actual labels. In logistic regression, we use a cost function called:

► Binary Cross-Entropy Loss

(Also known as Log Loss)

This is specifically designed for **binary classification problems**, where each output is either `0` or `1`.

Let's Start With Probabilities

Suppose:

- Your model predicts:
 $\hat{y} = P(y = 1 | x)$ — i.e., the probability that the output is 1 given input x
- The true label is $y \in \{0, 1\}$

We want our prediction \hat{y} to be **close to the actual label y** .

The Idea Behind Cross-Entropy

Cross-entropy comes from **information theory**, where it measures the distance between two probability distributions:

- One is the **true label y**
- The other is the **predicted probability \hat{y}**

Binary Cross-Entropy Formula:

$$\mathcal{L}(\hat{y}, y) = -[y \cdot \log(\hat{y}) + (1 - y) \cdot \log(1 - \hat{y})]$$

◆ Case 1: True label $y = 1$

Then:

$$\mathcal{L} = -[1 \cdot \log(\hat{y}) + 0 \cdot \log(1 - \hat{y})] = -\log(\hat{y})$$

If $\hat{y} = 0.9$, then loss = $-\log(0.9) = 0.105$ (small loss 

If $\hat{y} = 0.1$, then loss = $-\log(0.1) = 2.302$ (large loss 

◆ Case 2: True label $y = 0$

Then:

$$\mathcal{L} = -[0 \cdot \log(\hat{y}) + 1 \cdot \log(1 - \hat{y})] = -\log(1 - \hat{y})$$

If $\hat{y} = 0.1$, then loss = $-\log(0.9) = 0.105$ (good)

If $\hat{y} = 0.9$, then loss = $-\log(0.1) = 2.302$ (bad)

Key Insight:

- The loss is **small** when the model is confident and **correct**.
- The loss is **large** when the model is confident and **wrong**.

Cost Over All Training Samples

If you have m training examples:

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m \left[-y^{(i)} \log(\hat{y}^{(i)}) - (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \right]$$

Where:

- $\hat{y}^{(i)} = \sigma(w^T x^{(i)} + b)$ — prediction for example i
- $y^{(i)} \in \{0, 1\}$ — actual label

How does logistic regression works? (p2.2)



Why Logarithms in the Cost Function?

We use logarithms in the binary cross-entropy loss for two key reasons:

◆ 1. Penalize Wrong Predictions More

Logarithms grow rapidly negative as the predicted probability gets close to 0, which helps to strongly **punish confident wrong predictions**.

Example:

If the true label is 1:

- Predicting $\hat{y} = 0.9$:

$$\text{Loss} = -\log(0.9) \approx 0.105 \quad (\text{small penalty})$$

- Predicting $\hat{y} = 0.1$:

$$\text{Loss} = -\log(0.1) \approx 2.302 \quad (\text{huge penalty})$$

This ensures that the model learns to assign **high probability to the correct class**.

◆ 2. Natural Result of Maximum Likelihood Estimation (MLE)

If we assume that the target $y \in \{0, 1\}$ is a Bernoulli random variable, then the **likelihood** of a single training sample is:

$$P(y | x) = \hat{y}^y \cdot (1 - \hat{y})^{1-y}$$

To simplify optimization, we take the **log-likelihood**:

$$\log P(y | x) = y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})$$

We **negate** this because we want to **minimize** (instead of maximizing) the function:

$$\text{Loss} = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

This becomes the **binary cross-entropy loss function** used in logistic regression.

Reason for log function	Explanation
Punish wrong predictions	Log function penalizes low probabilities strongly
Derived from probability theory	It comes from the log of the Bernoulli likelihood in MLE
Smooth & differentiable	Great for optimization using gradient descent



$$h_{\theta}(x) = \frac{1}{1 + e^{-(w_0 + w_1x_1 + w_2x_2 + \dots + w_nx_n)}}$$

This is the **general logistic regression hypothesis** for multivariate data.

Decision Rule

$$\hat{y} = \begin{cases} 1 & \text{if } h_{\theta}(x) \geq 0.5 \\ 0 & \text{if } h_{\theta}(x) < 0.5 \end{cases}$$

Logistic regression uses **Binary Cross Entropy Loss**:

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

And optimized using **gradient descent**.